

Liver Disease Dataset

```
import numpy as np
import pandas as pd
import seaborn as sns
df = pd.read_csv("indian_liver_patient.csv")
df.head()

   Age  Gender  Total_Bilirubin  Direct_Bilirubin
Alkaline_Phosphotase \
0    65  Female             0.7                 0.1
187
1    62    Male              10.9                5.5
699
2    62    Male              7.3                 4.1
490
3    58    Male              1.0                 0.4
182
4    72    Male              3.9                 2.0
195

   Alamine_Aminotransferase  Aspartate_Aminotransferase
Total_Protiens \
0                      16                     18
6.8
1                      64                    100
7.5
2                      60                     68
7.0
3                      14                     20
6.8
4                      27                     59
7.3

   Albumin  Albumin_and_Globulin_Ratio  Dataset
0      3.3                  0.90          1
1      3.2                  0.74          1
2      3.3                  0.89          1
3      3.4                  1.00          1
4      2.4                  0.40          1

df

   Age  Gender  Total_Bilirubin  Direct_Bilirubin
Alkaline_Phosphotase \
0    65  Female             0.7                 0.1
187
```

1	62	Male	10.9	5.5
699				
2	62	Male	7.3	4.1
490				
3	58	Male	1.0	0.4
182				
4	72	Male	3.9	2.0
195				
...
...				
578	60	Male	0.5	0.1
500				
579	40	Male	0.6	0.1
98				
580	52	Male	0.8	0.2
245				
581	31	Male	1.3	0.5
184				
582	38	Male	1.0	0.3
216				
Alamine_Aminotransferase		Aspartate_Aminotransferase		
Total_Protiens \				
0		16		18
6.8				
1		64		100
7.5				
2		60		68
7.0				
3		14		20
6.8				
4		27		59
7.3				
...	
...				
578		20		34
5.9				
579		35		31
6.0				
580		48		49
6.4				
581		29		32
6.8				
582		21		24
7.3				
Albumin		Albumin_and_Globulin_Ratio	Dataset	
0	3.3		0.90	1
1	3.2		0.74	1

2	3.3	0.89	1
3	3.4	1.00	1
4	2.4	0.40	1
..
578	1.6	0.37	2
579	3.2	1.10	1
580	3.2	1.00	1
581	3.4	1.00	1
582	4.4	1.50	2

[583 rows x 11 columns]

df.tail(5)

	Age	Gender	Total_Bilirubin	Direct_Bilirubin
Alkaline_Phosphotase				\
578	60	Male	0.5	0.1
500				
579	40	Male	0.6	0.1
98				
580	52	Male	0.8	0.2
245				
581	31	Male	1.3	0.5
184				
582	38	Male	1.0	0.3
216				

	Alamine_Aminotransferase	Aspartate_Aminotransferase
Total_Protiens		\
578	20	34
5.9		
579	35	31
6.0		
580	48	49
6.4		
581	29	32
6.8		
582	21	24
7.3		

	Albumin	Albumin_and_Globulin_Ratio	Dataset
578	1.6	0.37	2
579	3.2	1.10	1
580	3.2	1.00	1
581	3.4	1.00	1
582	4.4	1.50	2

```
#Predict if a patient has liver disease (1) or not (0) using medical attributes
```

*****This is the following step on next

```
#Data Preprocessing  
#Exploring Dataset (Missing Values)  
#Encoding  
#Scalling  
#Balancing Classes  
#Feature Selection  
#Train Test Split
```

```
#Exploring Dataset (Missing Values)
```

```
df.isnull().sum()  
  
Age          0  
Gender       0  
Total_Bilirubin 0  
Direct_Bilirubin 0  
Alkaline_Phosphotase 0  
Alamine_Aminotransferase 0  
Aspartate_Aminotransferase 0  
Total_Protiens   0  
Albumin        0  
Albumin_and_Globulin_Ratio 4  
Dataset         0  
dtype: int64
```

```
df.dropna(inplace=True)
```

```
#df.dropna(inplace=True) Pandas ka ek statement hai jo missing values (NaN) wale rows ko remove karta hai.
```

```
df.shape
```

```
(579, 11)
```

```
#df.duplicated().sum() Pandas me duplicate rows count karne ke liye  
use hota hai.
```

```
df.duplicated().sum()
```

```
13
```

```
#df.drop_duplicates(inplace=True) Pandas ka use duplicate rows ko  
DataFrame se permanently remove karne ke liye hota hai.
```

```
df.drop_duplicates(inplace=True)
```

```
#Encoding
```

```
#Gender is a categorical feature with values Male and Female.
```

```
#Convert it into numeric form using Label Encoding or One-Hot Encoding
```

```
#Example male--1(Encoded_Value) female--0(Encoded_Value)
```

```
# Encoding helps ML models process non-numeric (categorical) data  
effectively.
```

```
#using sklearn
```

```
from sklearn.preprocessing import LabelEncoder, OrdinalEncoder
```

```
encoder=LabelEncoder()
```

```
df['Gender']=encoder.fit_transform(df['Gender'])
```

```
df.head(2)
```

```
   Age  Gender  Total_Bilirubin  Direct_Bilirubin  
Alkaline_Phosphotase \
0    65        0             0.7            0.1
187
1    62        0            10.9            5.5
699
```

```
   Alamine_Aminotransferase  Aspartate_Aminotransferase  
Total_Protiens \
0                      16                  18
6.8
1                      64                 100
7.5
```

```
   Albumin  Albumin_and_Globulin_Ratio  Dataset
0      3.3                  0.90          1
1      3.2                  0.74          1
```

```
df.tail
```

	Direct_Bilirubin	Alkaline_Phosphotase	Age	Gender	Total_Bilirubin
0	65	0	0.7	\	0.1
187					
1	62	0	10.9		5.5
699					
2	62	0	7.3		4.1
490					
3	58	0	1.0		0.4
182					
4	72	0	3.9		2.0
195					
..
...					
578	60	0	0.5		0.1
500					
579	40	0	0.6		0.1
98					
580	52	0	0.8		0.2
245					
581	31	0	1.3		0.5
184					
582	38	0	1.0		0.3
216					
	Alamine_Aminotransferase	Aspartate_Aminotransferase			
Total_Protiens	\				
0		16			18
6.8					
1		64			100
7.5					
2		60			68
7.0					
3		14			20
6.8					
4		27			59
7.3					
..
...					
578		20			34
5.9					
579		35			31
6.0					
580		48			49
6.4					
581		29			32
6.8					
582		21			24
7.3					

```

      Albumin  Albumin_and_Globulin_Ratio  Dataset
0          3.3                      0.90      1
1          3.2                      0.74      1
2          3.3                      0.89      1
3          3.4                      1.00      1
4          2.4                      0.40      1
..          ...
578         1.6                      0.37      2
579         3.2                      1.10      1
580         3.2                      1.00      1
581         3.4                      1.00      1
582         4.4                      1.50      2

```

[566 rows x 11 columns]>

```

# manually
# df['Gender'] = df['Gender'].map({"Male":1, "Female":0})

```

⌚ Feature Scaling

**Scaling ensures all numerical features are on a similar range for better model performance.

**Common methods: StandardScaler: Scales data to mean = 0 and std = 1 MinMaxScaler: Scales values between 0 and 1

Example:

Feature Before After (MinMax)

Age 65 0.78

Total_Bilirubin 10.9 0.85

```
# Apply scaling to all numeric columns (except target) before model training.
```

```
from sklearn.preprocessing import StandardScaler ,MinMaxScaler
```

```
features = ['Age', 'Total_Bilirubin', 'Direct_Bilirubin',
           'Alkaline_Phosphotase', 'Alamine_Aminotransferase',
           'Aspartate_Aminotransferase', 'Total_Protiens', 'Albumin',
           'Albumin_and_Globulin_Ratio']
```

```

scaler=StandardScaler()
df[features]=scaler.fit_transform(df[features])
df.head(5)

      Age  Gender  Total_Bilirubin  Direct_Bilirubin
Alkaline_Phosphotase \
0  1.236928      0       -0.420124       -0.495190
0.429625
1  1.052432      0        1.203777       1.406906
1.654054
2  1.052432      0        0.630636       0.913770
0.803490
3  0.806437      0       -0.372362       -0.389518
0.449974
4  1.667418      0        0.089335       0.174066
0.397068

      Alamine_Aminotransferase  Aspartate_Aminotransferase
Total_Protiens \
0                  -0.352659       -0.315148
0.280819
1                  -0.088755       -0.033926
0.925059
2                  -0.110747       -0.143671
0.464887
3                  -0.363655       -0.308289
0.280819
4                  -0.292181       -0.174537
0.740991

      Albumin  Albumin_and_Globulin_Ratio  Dataset
0  0.194225           -0.150315         1
1  0.068445           -0.651328         1
2  0.194225           -0.181628         1
3  0.320004           0.162818         1
4 -0.937791           -1.715981         1

```

#Data Balancing

Class Balance

```
* Check if target classes (0 = No Disease, 1 = Liver Disease) are balanced.
```

```

* Imbalanced data can bias the model toward the majority class.

* Use techniques like:
  -SMOTE / Oversampling → add minority samples

  -Undersampling → reduce majority samples

!pip install -q imbalanced-learn

from imblearn.over_sampling import SMOTE

X = df.drop('Dataset',axis=1)
y = df['Dataset']

smote = SMOTE(random_state=42)

X_resampled, y_resampled = smote.fit_resample(X, y)

y_resampled.value_counts()

Dataset
1    404
2    404
Name: count, dtype: int64

#Now i move Train -Test Split

from sklearn.model_selection import train_test_split

x_train,x_test,y_train,y_test=train_test_split(X_resampled,y_resampled
,test_size=0.2,random_state=42)

```

#Train and Eval Models

logistic regression -SVC -Decision Tree -Random Forest

```

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix,classification_report

model=LogisticRegression()
model.fit(x_train,y_train)

y_pred=model.predict(x_test)

```

```
print("\n\n",classification_report(y_test, y_pred))
print("\n\n", confusion_matrix(y_test, y_pred))
```

	precision	recall	f1-score	support
1	0.80	0.57	0.66	76
2	0.69	0.87	0.77	86
accuracy			0.73	162
macro avg	0.75	0.72	0.72	162
weighted avg	0.74	0.73	0.72	162

```
[[43 33]
 [11 75]]
```

#Random Forest

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix,classification_report

model_r=RandomForestClassifier()
model_r.fit(x_train,y_train)

y_pred=model_r.predict(x_test)

print("\n\n",classification_report(y_test, y_pred))
print("\n\n", confusion_matrix(y_test, y_pred))
```

	precision	recall	f1-score	support
1	0.83	0.76	0.79	76
2	0.80	0.86	0.83	86
accuracy			0.81	162
macro avg	0.82	0.81	0.81	162
weighted avg	0.82	0.81	0.81	162

```
[[58 18]
 [12 74]]
```

```
#SVC

from sklearn.svm import SVC
from sklearn.metrics import confusion_matrix, classification_report

model = SVC()

model.fit(x_train,y_train)

y_pred = model.predict(x_test)

print("\n\n",classification_report(y_test, y_pred))

print("\n\n", confusion_matrix(y_test, y_pred))
```

	precision	recall	f1-score	support
1	0.84	0.61	0.70	76
2	0.72	0.90	0.80	86
accuracy			0.76	162
macro avg	0.78	0.75	0.75	162
weighted avg	0.77	0.76	0.75	162

```
[[46 30]
 [ 9 77]]
```

#now i find best accuracy using Random Forest

```
# test 1
import numpy as np

pred = model_r.predict(np.array([df.iloc[0,:-1]]))

if pred[0] == 1:
    print("Liver Disease")
else:
    print("No Liver Disease")
```

Liver Disease

```
D:\narayan\Python\new\Lib\site-packages\sklearn\base.py:493:
UserWarning: X does not have valid feature names, but
RandomForestClassifier was fitted with feature names
warnings.warn(
```

```
# test 2

import numpy as np

pred = model_r.predict(np.array([df.iloc[23,:-1]]))

if pred[0] == 1:
    print("Liver Disease")
else:
    print("No Liver Disease")

No Liver Disease

D:\narayan\Python\new\Lib\site-packages\sklearn\base.py:493:
UserWarning: X does not have valid feature names, but
RandomForestClassifier was fitted with feature names
    warnings.warn(
#Test

actual = df.iloc[0, -1]
print("Actual Value:", actual)

Actual Value: 1

x_test_row = df.iloc[[0], :-1]  # double brackets
pred = model_r.predict(x_test_row)

print("Predicted Value:", pred[0])

Predicted Value: 1

if pred[0] == actual:
    print(" Model prediction is CORRECT")
else:
    print(" Model prediction is WRONG")

 Model prediction is CORRECT

proba = model_r.predict_proba(X_test_row)
print("Prediction Probabilities:", proba)

Prediction Probabilities: [[0.87 0.13]]
```